



MetroMetric

Can We Improve Transit Through
Data Science?
Austin Brown ~ 3/16/2014 ~ DAT4

Buses Are Great



(But a pain to use)

<http://www.howwedrive.com/2010/06/28/german-efficiency/>

Münster

Background

- Bus system economics, congestion reduction, and energy benefits depend directly on ridership
- All buses in Washington Metro Area Transit Authority (WMATA) are tracked by GPS, with information available by API
- Bus arrival predictions are available real-time through the website and another API
- Good bus predictions are likely to improve ridership and user satisfaction
- Bad bus predictions are certain to frustrate potential riders

Open Transit Data

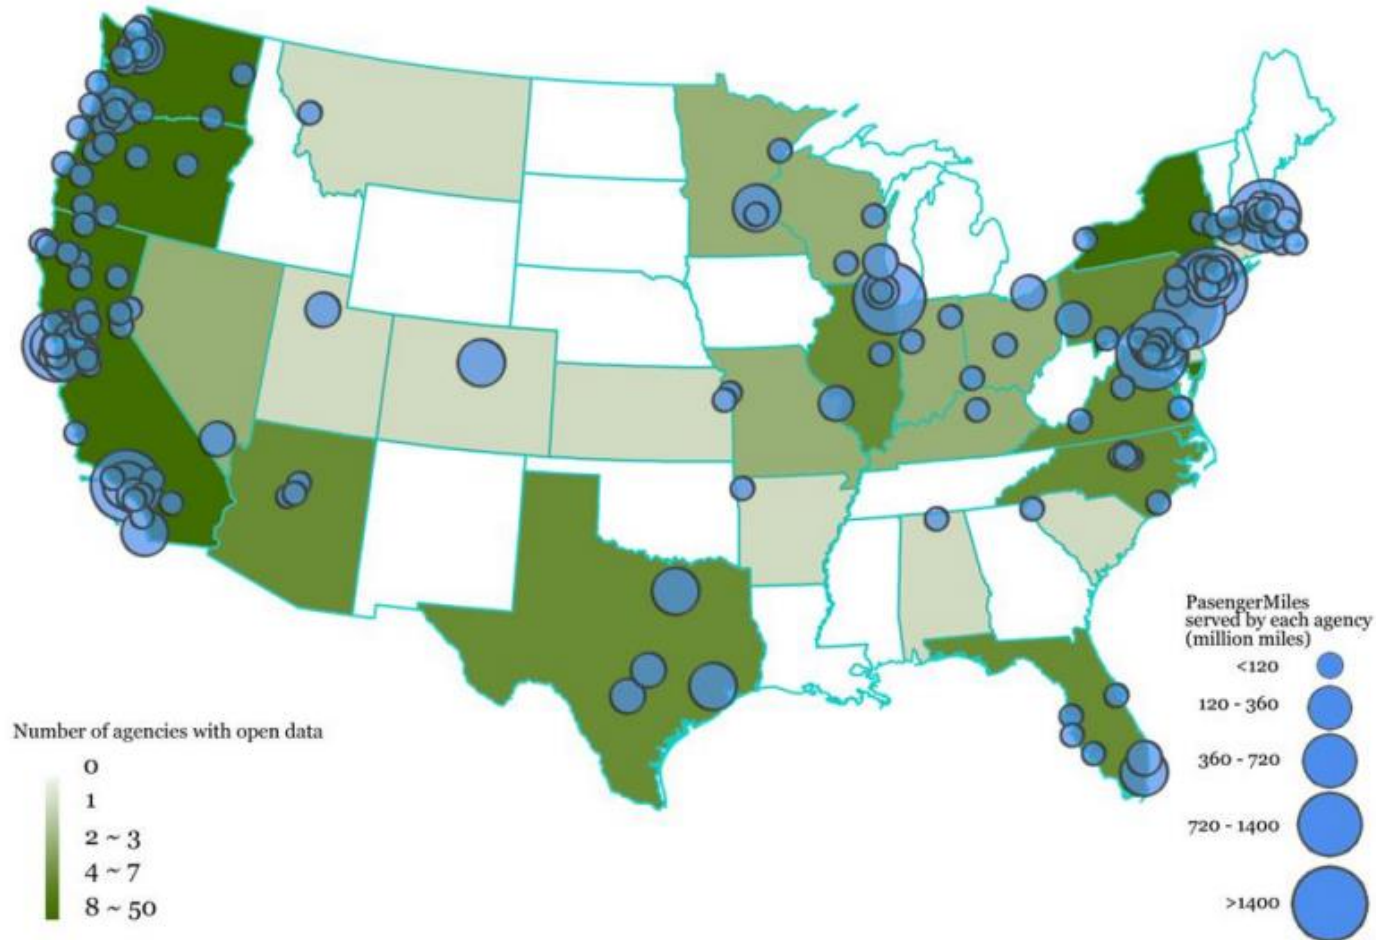
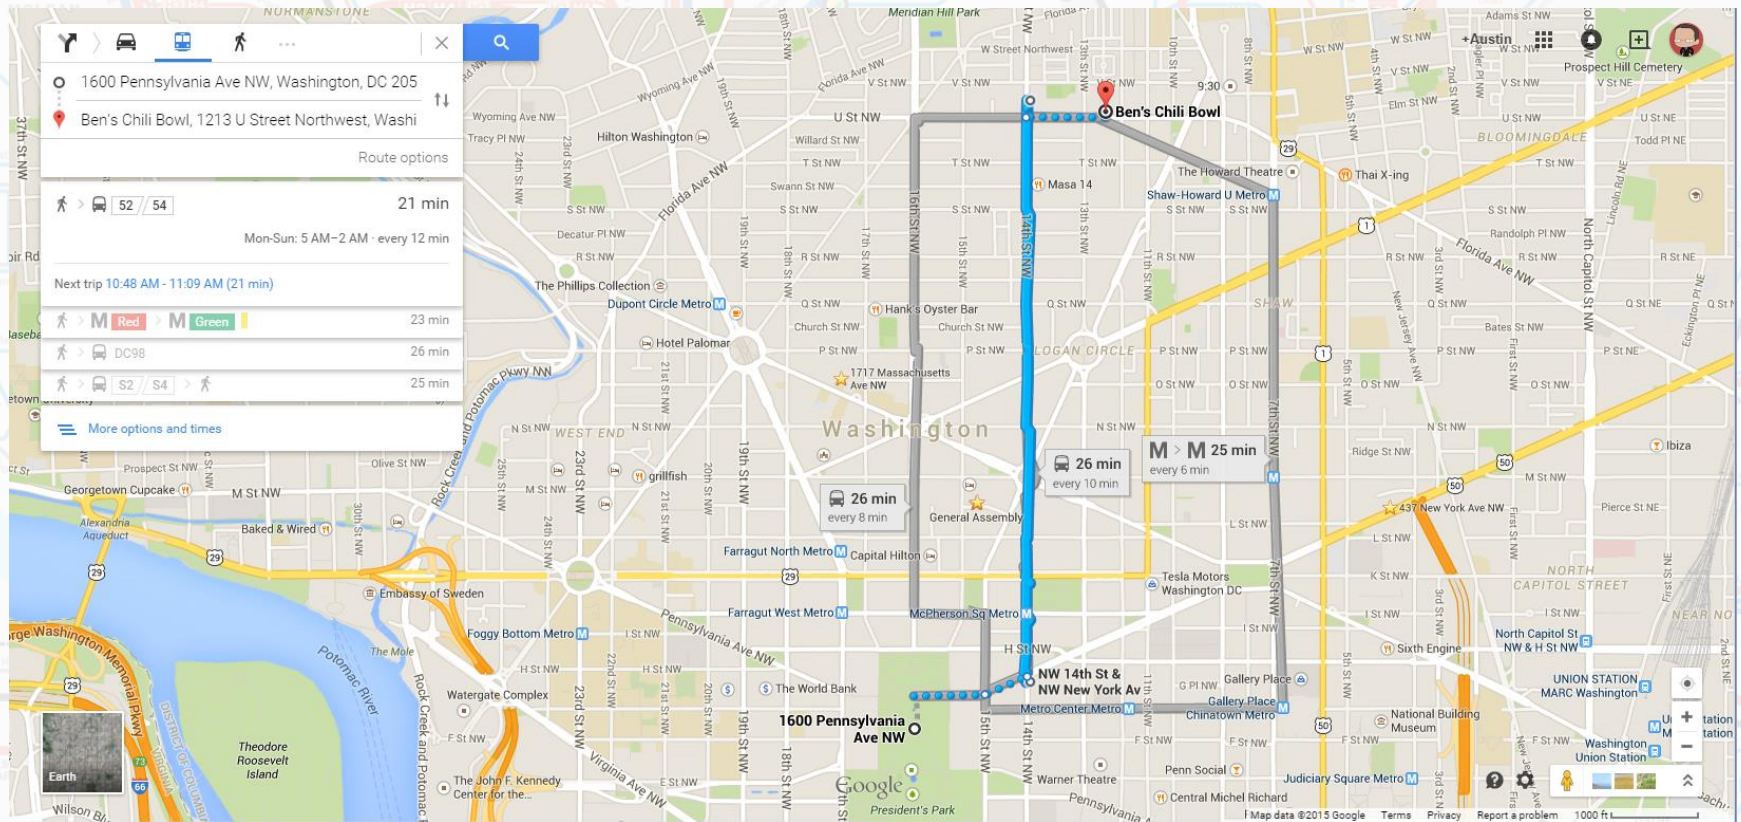


Figure 1- Number and size of agencies with open data by state.

Illustration from <http://www.prism.gatech.edu/~lreed3/opendata.pdf>

Example Application



Example Application



Personal Example

7:01 AM

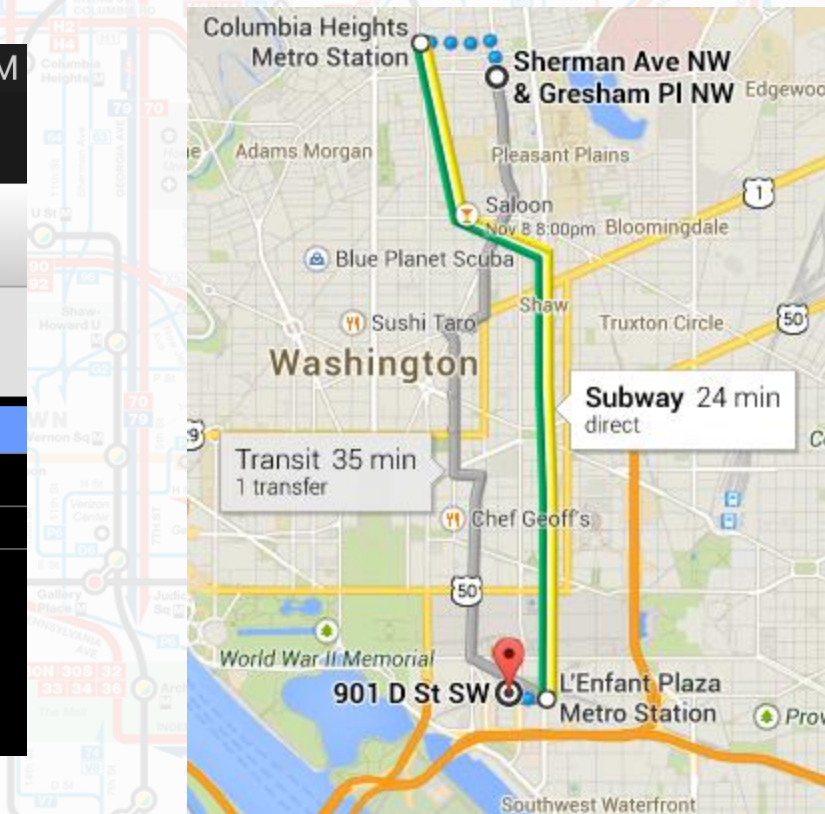
Bus Results

Home Refresh Favorites

Route: 63 TAKOMA-PETWORTH
Direction: South to Federal Triangle
Stop: Sherman Ave + Girard St (#1001930)

Predictions for this route

4 Minutes
26 Minutes
36 Minutes



DC Metro Transit app

Google Maps

When NextBus Goes Wrong

- A bad prediction can be worse than no prediction
- Especially early on, next bus was littered with “ghost” buses, and predictions seemed problematic
- The NextBus algorithm is proprietary so the data community can’t work on improving it directly

Question

- How good are NextBus predictions?
- Can we mash up NextBus with other data to improve predictions?


Data Structure Concept

Field(s)	Description	Source
Route, Direction	e.g. 70	Route API
StopID, Lat, Lon	Stop Identified and location	Route API
BusID, Lat, Lon	Bus Identified and location	BusLocation API
Deviation	How far behind the bus says it is	BusLocation API
TripID	Unique trip identifier	BusLocation / Nextbus API
Time / Date	Time the prediction was made	Datetime package
Weather / Temp	Reported weather	Openweathermap API
Predicted Arrival	In minutes, when the bus should arrive	NextBus API
Actual Arrival	In minutes, how long from the prediction the bus arrived	Calculated

The basic data unit is one prediction. Data collection is done from four APIs in real time and stored in a data frame / csv

WMATA API Key

Your subscriptions

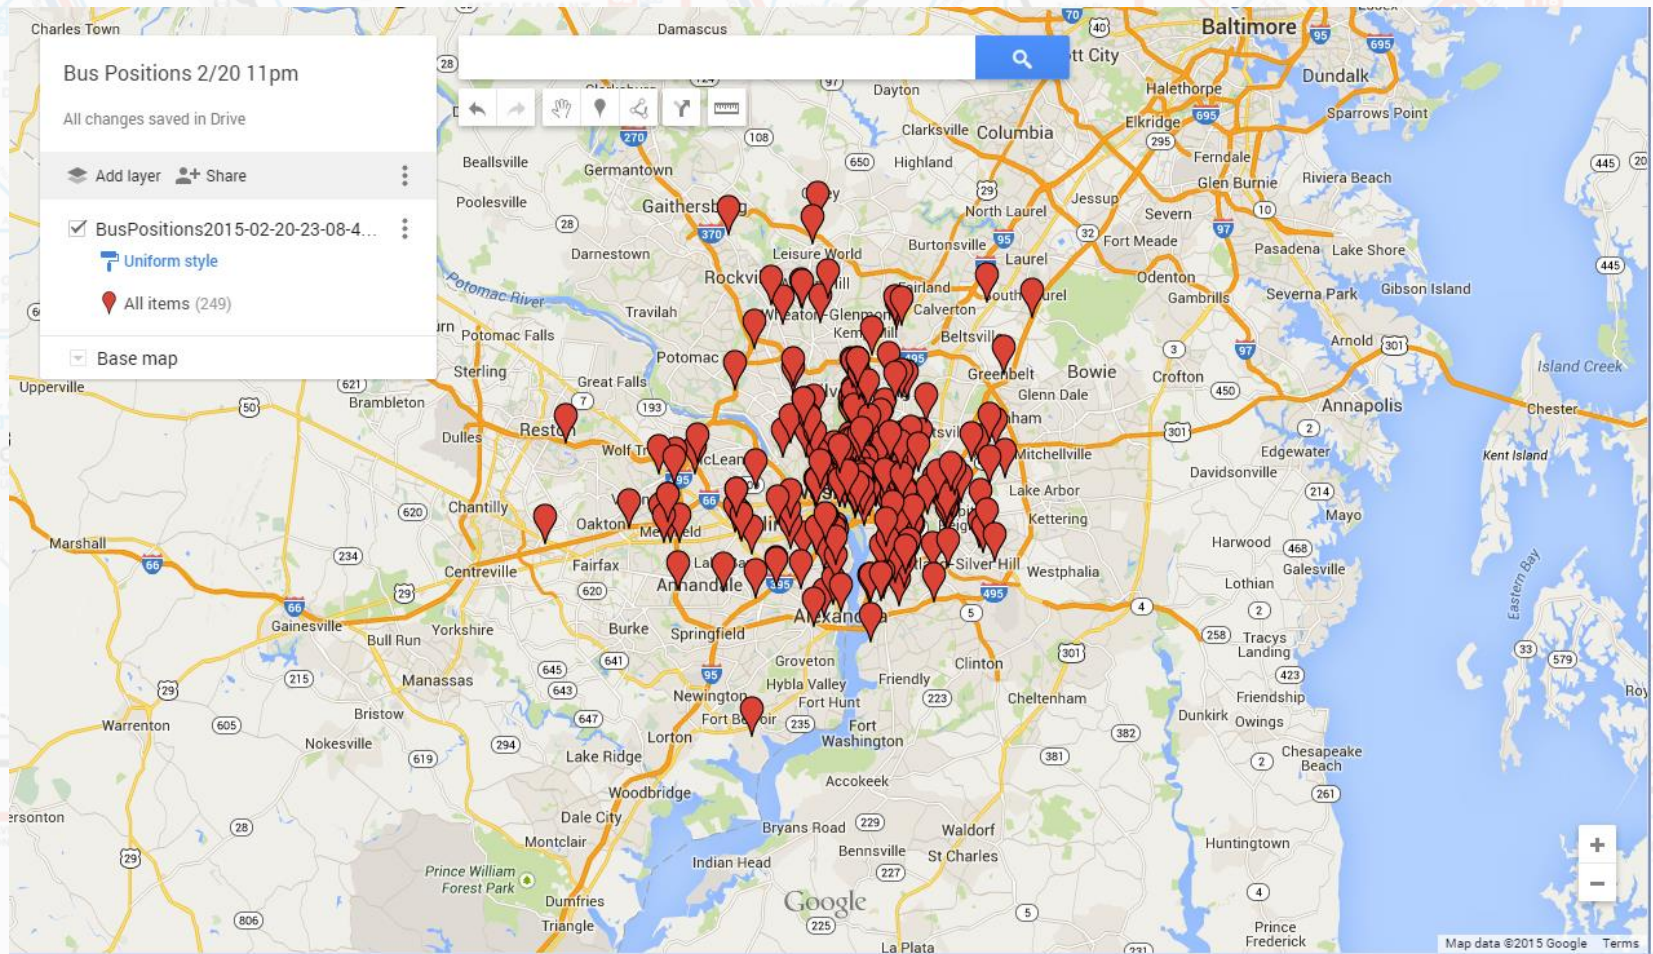
 Analytics reports

Subscription details

			Product	State
Subscription name:	Legacy Subscription - qkk2x7wckemx2bxgs8g2sq8a	Rename	Default Tier	Active ✕ Cancel
Started on:	12/11/2014			
Primary key:	XXXXXXXXXXXXXXXXXXXXXXXXXXXX	Show Regenerate		
Secondary key:	XXXXXXXXXXXXXXXXXXXXXXXXXXXX	Show Regenerate		
Subscription name:	Temporary Key - Please do not break the API	Rename	Tier 3	Active ✕ Cancel
Started on:	2/12/2015			
Primary key:	XXXXXXXXXXXXXXXXXXXXXXXXXXXX	Show Regenerate		
Secondary key:	XXXXXXXXXXXXXXXXXXXXXXXXXXXX	Show Regenerate		

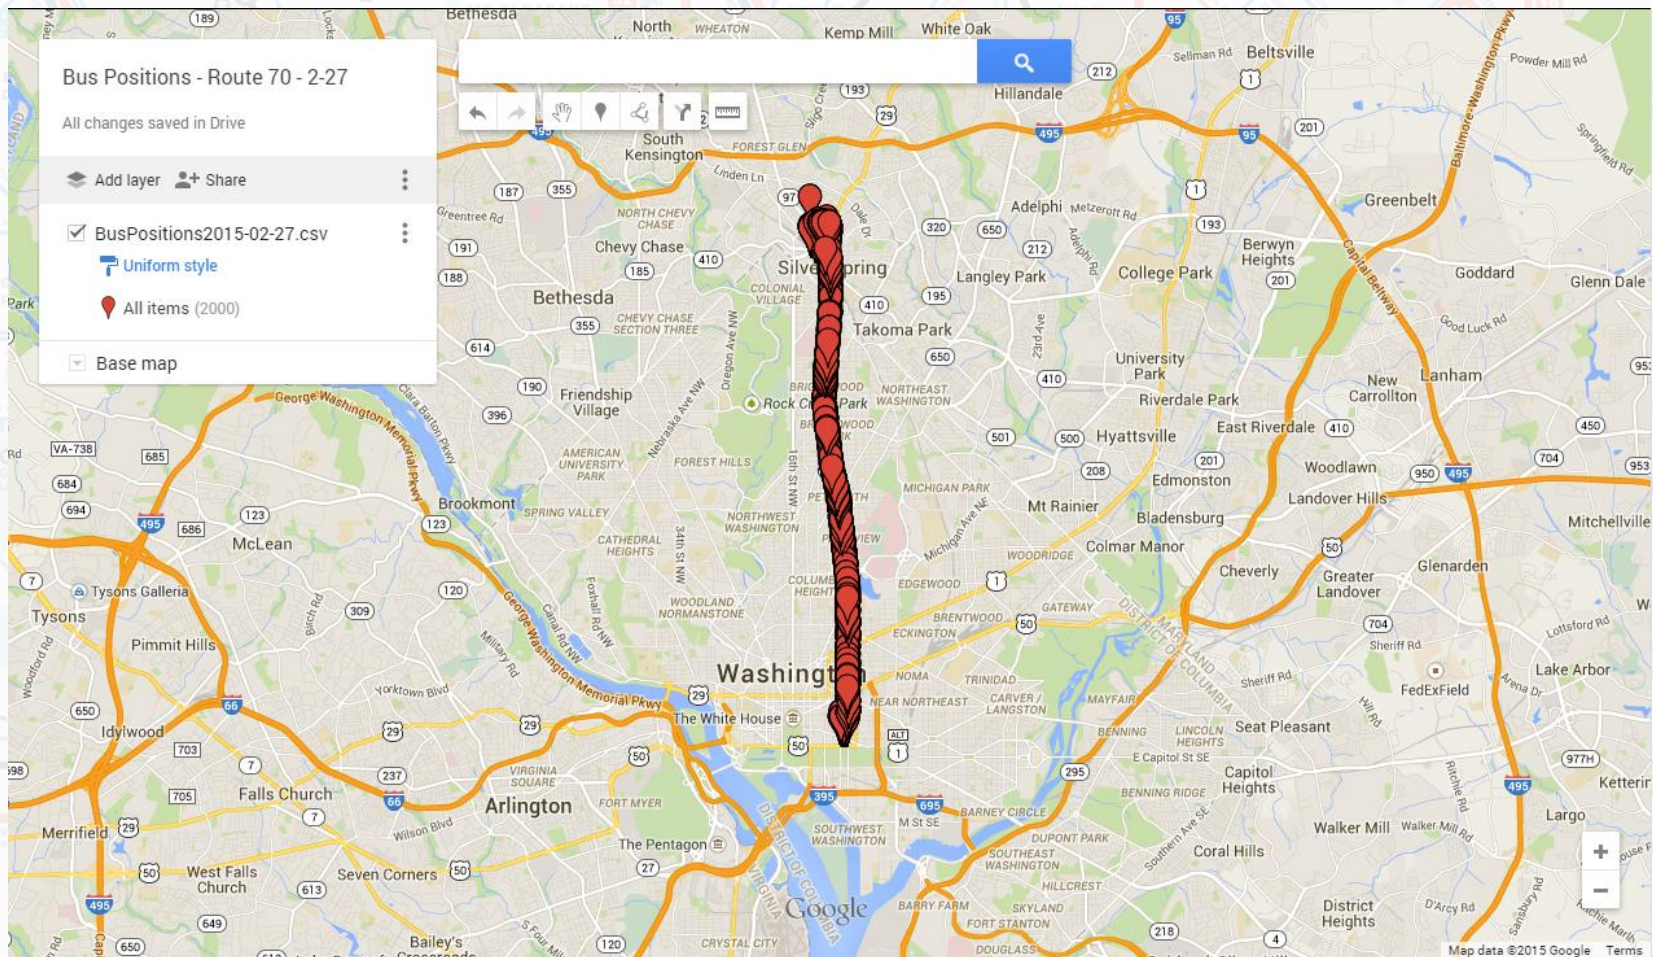
I need a lot of calls... WMATA to the rescue

WMATA APIs – Bus Positions



One data call for all bus positions. 249 buses in system (at 11 pm)

WMATA APIs – Bus Positions



Many data calls for route 70 only

Data Collection Overview

- Days collected:
 - 3/5, snow day, 5 am to 3 pm, with some data gaps
 - 3/8, Saturday, 8:30 pm to midnight
 - 3/9, Sunday, 5 am to 10:45
 - 3/10, Monday, 5 am to 8:30 am
 - 3/12, 7:30 am to 5:45 pm
- Frame Size
 - 500,000 future predictions, 10,000 arrivals
- ~100 MB on disk as csv

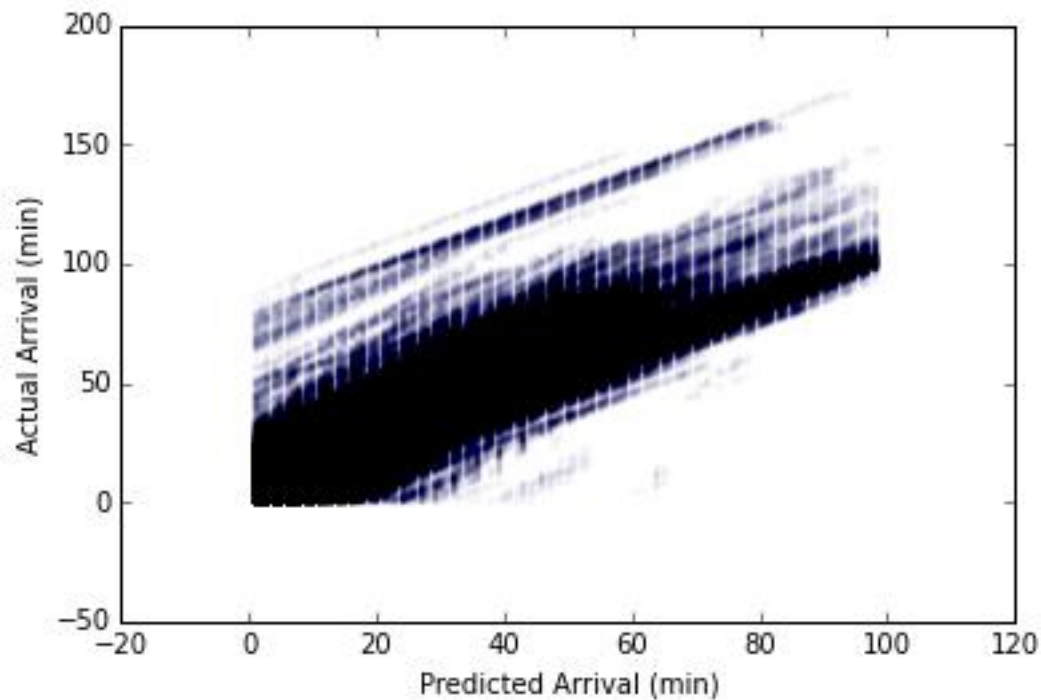
Method: Estimating Arrivals

- After I have a big database of predictions, identify arrivals
- 2 possible methods
 - Predicted arrival = 0 minutes
 - Lat/Lon check
- Pull the “predictions” where the bus has arrived and join that data frame to the predictions to assign the correct arrival time for each prediction
 - In my test data, covers 99% of all predictions
- The ‘Error’ is the difference between predicted arrival and actual arrival

A “Quality” Metric For Predictions


- Attempted to be reflective of our expectations for a prediction:
 - <1 is “perfect”, far off = sucky.
 - Should be more forgiving for further outCoub be more critical of being wrong early (since you will miss the bus?)
- For this pass, I used $[1 - |\text{predicted} - \text{actual}| / \text{actual}]$
 - Predicted 2 minutes, actual 4 minutes = 0.5
 - Predicted 5 minutes, actual 7 minutes = 0.71
 - Predicted 10 minutes, actual 12 minutes = 0.83

Results: NextBus Performance



Results: NextBus Performance


All Predictions



Actual Arrival – Predicted Arrival (positive means the bus was later than predicted)


Results: NextBus Performance

Only predictions < 10 mins



Actual Arrival – Predicted Arrival (positive means the bus was later than predicted)

“Quality” Metric




Quality

mean	0.641955
std	3.706173
min	-1198.000000
25%	0.647815
50%	0.782870
75%	0.893502
max	1.000000

Recall 0.5 means it was as far off as the total time the bus took to arrive.
About half of predictions are better than 0.8, meaning ~ within 20%


Results: NextBus Performance

- Some evidence (from one day!) that snow increases error




Results: NextBus Performance

- Lateness looks like somewhat of a predictor of more lateness



Results: NextBus Performance





Regression To Improve Predictions

- Variables considered
 - Predicted Arrival (hopefully this will be the major predictor!)
 - Deviation (how late the bus is reporting itself to be)
 - Weather (text descriptor, created dummy variable)
 - Time (for now, as a linear variable... non-ideal)

Regression Results



- On all data:
 - Predicted arrival: 1.107 (predictions high by 10% on average)
 - Deviation: -0.0007
 - BadWeather: 1.028 (on average rain / snow increase error)
 - Time: 0.0002
 - $R^2 = 0.836$ (same when done w/ train-test-split)

Regression Results



- Not really an improvement, seems to fix one problem and cause another

Regression Results



- On quality metric, looks better – but I think this means my quality metric isn't great...

Improvement: Training Variables

- Try:
 - Time of day dummy variables (rush hour?)
 - Day of week
 - Bus ID (as a proxy for driver)
 - Temperature (another proxy for weather?)
 - Location / stop ID (downtown?)

Other Next Steps





























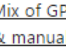

- More complete data collection
 - More routes
 - All stops
 - Consistent time coverage
- Use WMATA Incident API to explore effect of reported incidents on bus arrival predictions
- Better optimization functions / other algorithms besides linear regression
 - e.g. Lasso to reduce useless features, or random forests to divide by classes (rush hour, weather, etc)
- Explore outlier buses – is there an algorithm to identify errors in real time?
- Fun stuff like animated maps

Thank You!



DC Transit Agencies Status (2012)

	Schedule data			Real-time data	
	Public GTFS	Shapes in GTFS	On Google	Tracking	Tracking API
Metrorail	 Here		 and Bing		 Custom
Metrobus		Most ¹	 and Bing		 Custom
Circulator (DC)	 WMATA ²		 WMATA ²		 Nextbus
ART (Arlington)	 Here		 In process		 Connexionz
DASH (Alexandria)	Via email only ³				
Ride On (Montgomery)	 Old ⁴				 More info
The Bus (Prince George's)					 Nextbus
MTA (Maryland) commuter bus	 Here				
MARC	 Confusingly ⁵			 Here	

	Schedule data			Real-time data	
	Public GTFS	Shapes in GTFS	On Google	Tracking	Tracking API
MARC	 Confusingly ⁵			 Here	
Fairfax (County) Connector					
CUE (Fairfax City)					 Nextbus
Loudoun County Transit				 Text/email alerts	
PRTC					
VRE	 Unofficial ⁶			 Mix of GPS & manual ⁷	

¹ WMATA's GTFS file contains most Metrobus routes, but some paths cut diagonally across the grid over some long sections such as freeway or bridge segments of routes.

² Circulator route and schedule data is included as part of the WMATA GTFS feed. However, there are some quality issues such as route names.

³ DASH feed is not publicly available, but officials can provide it via email.

⁴ Ride On's feed no longer appears to be on their website. [GTFS Data Exchange](#) has cached a version from December 2010 which was apparently posted in a news release.

⁵ MARC lines are listed in the MTA Maryland feed as lines 300, 301, and 302, which doesn't very easily differentiate them for someone unfamiliar with their GTFS feed.

⁶ Someone not affiliated with VRE created a GTFS file in 2009, but it hasn't been updated since and VRE does not offer an official one.

⁷ VRE has a [page with train status](#) which lists some trains' positions through GPS and some from manual reports from the conductor.

Notes

- GTFS: General Transit Feed Specification (originally Google)
- GTFS feed is composed of a series of text files collected in a ZIP file. Each file models a particular aspect of transit information: stops, routes, trips, and other schedule data.
- Routes can include shapeform coordinates

APIs Used

- WMATA Route Structure
 - To get list and lat / lon of stops
- WMATA Bus Position
 - To get list and lat / lon of buses, as well as 'deviation'
- WMATA NextBus Predictions
 - To get predicted arrivals for each stop
- OpenWeatherMap